


Lecture 3

1. Newton's method
2. Exponential family
3. Generalized Linear Models (GLM)
4. Generative learning algorithm

1 Newton's method

Newton's method is an algorithm that help to solve an equation $f(\theta) = 0$.



We start from some random $\theta^{(0)}$, by definition:

$$f'(\theta^{(0)}) = \frac{f(\theta^{(0)})}{\Delta},$$

which implies that $\Delta = \frac{f(\theta^{(0)})}{f'(\theta^{(0)})}$. Then

$$\theta^{(1)} = \theta^{(0)} - \frac{f(\theta^{(0)})}{f'(\theta^{(0)})}$$

or, in general,

$$\theta^{(t+1)} = \theta^{(t)} - \frac{f(\theta^{(t)})}{f'(\theta^{(t)})}.$$

To maximize likelihood $l(\theta)$ we find θ such that $l'(\theta) = 0$, then one iteration of the Newton's methods is written as

$$\theta^{(t+1)} = \theta^{(t)} - \frac{l'(\theta^{(t)})}{l''(\theta^{(t)})}.$$

This algorithm has a quadratic convergence. It means that error decreasing as square after each iteration. For example, in linear regression we need just 3 iterations to find the parameters with very good accuracy.

When the problem is to find several parameters, one iteration of the Newton's methods can be written as a vector equation:

$$\theta^{(t+1)} = \theta^{(t)} - H^{-1} \nabla_{\theta} l,$$

where θ is a vector and H is the Hessian matrix:

$$H_{ij} = \frac{\partial^2 l}{\partial \theta_i \partial \theta_j}.$$

Notice that if we want to minimize something (instead of maximization) the algorithm does not change.

2 Exponential family

In this section we generalise the ideas of linear and logistic regressions. Before we considered two cases:

- $y \in \mathbb{R} \Rightarrow$ Gaussian distribution $N(\mu, \sigma^2) \Rightarrow$ linear regression
- $y \in \{0, 1\} \Rightarrow$ Bernoulli distribution with parameter φ such that $p(y = 1 | \varphi) = \varphi \Rightarrow$ logistic regression

We will show that in both cases we deal with particular cases of general class of exponential family distributions:

$$p(y; \eta) = b(y) \exp(\eta^T T(y) - a(\eta)),$$

where η is a natural parameter, $T(y)$ is a sufficient statistics (in many cases $T(y) = y$).

For different choices of a , b and T we will have different distributions:

1. Bernoulli distribution with φ : $p(y = 1; \varphi) = \varphi$. Then

$$\begin{aligned} p(y; \varphi) &= \varphi^y (1 - \varphi)^{1-y} = \exp(\ln(\varphi^y (1 - \varphi)^{1-y})) = \\ &= \exp(y \ln \varphi + (1 - y) \ln(1 - \varphi)) = \exp\left(\ln \frac{\varphi}{1 - \varphi} y + \ln(1 - \varphi)\right). \end{aligned}$$

By introducing the notations $\eta = \ln \frac{\varphi}{1 - \varphi}$, $T(y) = y$, $-a(\eta) = \ln(1 - \varphi)$, $b(y) = 1$, we obtain the exponential family formula. Notice that we can find

$$\varphi = \frac{1}{1 + e^{-\eta}},$$

then $a(\eta) = \ln(1 + e^{\eta})$.

2. Gaussian distribution $N(\mu, \sigma^2)$. For simplicity we set $\sigma^2 = 1$. Then

$$\begin{aligned} p(y; \mu) &= \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}(y - \mu)^2\right) = \dots = \\ &= \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}y^2\right) \exp\left(\mu y - \frac{1}{2}\mu^2\right). \end{aligned}$$

By introducing the notations $b(y) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}y^2\right)$, $\eta = \mu$, $T(y) = y$, $a(\eta) = \frac{1}{2}\eta^2$, we obtain the exponential family formula.

3 Generalized Linear Models (GLM)

The Generalized Linear Models use the exponential family distributions to build different algorithms in one general framework. Assuming some distribution for the output y we obtain the form for the hypothesis $h_\theta(x)$ and find the parameters by finding the maximum of log-likelihood. We assume the following:

- $y \mid x; \theta \sim \text{ExpFamily}(\eta)$
- given x our goal is to output $E[T(y) \mid x]$, in other words,

$$h(x) = E[T(y) \mid x]$$

- $\eta = \theta^T x$ (in more general case if $\eta \in \mathbb{R}^k$, then $\eta_i = \theta_i^T x$)

Examples of GLM:

1. Bernoulli distribution

- $y \mid x; \theta \sim \text{ExpFamily}(\eta)$
- for fixed x and θ , the algorithm output is

$$h_\theta(x) = E[y \mid x; \theta] = p(y = 1 \mid x; \theta) = \varphi = \frac{1}{1 + e^{-\eta}} = \frac{1}{1 + e^{-\theta^T x}}$$

Definition.


$$g(\eta) = E[y \mid \eta] = \frac{1}{1 + e^{-\eta}} \text{ is a canonical response function}$$

g^{-1} is a canonical link function

2. Gaussian distribution

Exercise for the homework

3. Multinomial distribution: $y \in \{1, \dots, k\}$



Parameters $\varphi_1, \varphi_2, \dots, \varphi_k$ are defined as $p(y = i) = \varphi_i$. Such formulation is redundant, because φ_k can be expressed as $\varphi_k = 1 - \varphi_1 - \dots - \varphi_{k-1}$. That is why we do not take φ_k as a parameter. We introduce function $T(y), y \in \{1, \dots, k\}$ as follows:

$$T(1) = \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, T(2) = \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \dots, T(k-1) = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}$$

and

$$T(k) = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, T(1), T(2), \dots, T(k) \in \mathbb{R}^{k-1}.$$

The indicator function $\mathbb{1}\{True\} = 1, \mathbb{1}\{False\} = 0$, for example, $\mathbb{1}\{2 = 3\} = 0$. Then $T(y)_i = \mathbb{1}\{y = i\}$. The distribution:

$$\begin{aligned} P(y) &= \varphi_1^{\mathbb{1}\{y=1\}} \cdot \varphi_2^{\mathbb{1}\{y=2\}} \cdot \dots \cdot \varphi_k^{\mathbb{1}\{y=k\}} = \\ &= \varphi_1^{T(y)_1} \cdot \varphi_2^{T(y)_2} \cdot \dots \cdot \varphi_{k-1}^{T(y)_{k-1}} \cdot \varphi_k^{1 - \sum_{j=1}^{k-1} T(y)_j} = \dots = \\ &= b(y) \exp(\eta^T T(y) - a(\eta)), \end{aligned}$$

where

$$\eta = \begin{bmatrix} \ln(\varphi_1/\varphi_k) \\ \vdots \\ \ln(\varphi_{k-1}/\varphi_k) \end{bmatrix} \in \mathbb{R}^{k-1}, a(\eta) = -\ln \varphi_k, b(y) = 1.$$

We could solve these equations with respect to η and obtain

$$\varphi_i = \frac{e^{\eta_i}}{1 + \sum_{j=1}^{k-1} e^{\eta_j}} = \frac{e^{\theta_i^T x}}{1 + \sum_{j=1}^{k-1} e^{\theta_j^T x}}, i = 1 \dots k-1.$$

The purpose of these transformation is to obtain the learning algorithm:

$$h_{\theta}(x) = E[T(y) | x; \theta] = E \left[\begin{array}{c} \mathbb{1}\{y = 1\} \\ \vdots \\ \mathbb{1}\{y = k - 1\} \end{array} \middle| x; \theta \right] =$$

$$= \begin{bmatrix} \varphi_1 \\ \vdots \\ \varphi_{k-1} \end{bmatrix} = \begin{bmatrix} \frac{e^{\theta_1^T x}}{1 + \sum_{j=1}^{k-1} e^{\theta_j^T x}} \\ \vdots \\ \frac{e^{\theta_{k-1}^T x}}{1 + \sum_{j=1}^{k-1} e^{\theta_j^T x}} \end{bmatrix}$$

This algorithm is called **softmax regression** - generalisation of logistic regression for k classes. To fit the parameters for the given training set $(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})$, we find the likelihood in the form

$$L(\theta) = \prod_{i=1}^m p(y^{(i)} | x^{(i)}; \theta) = \prod_{i=1}^m \varphi_1^{\mathbb{1}\{y^{(i)}=1\}} \cdot \varphi_2^{\mathbb{1}\{y^{(i)}=2\}} \cdot \dots \cdot \varphi_k^{\mathbb{1}\{y^{(i)}=k\}},$$

where $\varphi_1 = \frac{e^{\theta_1^T x}}{1 + \sum_{j=1}^{k-1} e^{\theta_j^T x}}$. The last step will be to find the log-likelihood, find the

derivatives and use stochastic gradient descent to maximize the log-likelihood function with respect to parameters θ .

3.1 Python implementation

Loading necessary libraries:

```
In [1]: from matplotlib import cm
import matplotlib.pyplot as plt
import numpy as np
import sklearn.datasets as ds
```

The softmax regression in `scikit-learn` can be trained using `LogisticRegression` class. We generate data with 3 classes and 2 features first:


```
In [2]: from sklearn.linear_model import LogisticRegression
```

```
In [3]: X,y = ds.make_classification(n_features=2,
                                   n_redundant=0,
                                   n_informative=2,
                                   n_clusters_per_class=1,
```

```

n_classes=3,
n_samples=200,
random_state=3216)
ix0 = [i for i,x in enumerate(y) if x == 0]
ix1 = [i for i,x in enumerate(y) if x == 1]
ix2 = [i for i,x in enumerate(y) if x == 2]
fig = plt.figure(figsize=(6,6))
#ax = plt.axes(xlim=(-3.1, 3.1), ylim=(-3.1, 3.1))
plt.scatter(X[ix0,0],X[ix0, 1],marker='o',color='red')
plt.scatter(X[ix1,0],X[ix1, 1],marker='x',color='blue')
plt.scatter(X[ix2,0],X[ix2, 1],marker='s',color='green')
plt.show()

```



Training the softmax regression is easy as before:

```

In [4]: model = LogisticRegression(multi_class='multinomial', solver='newton-cg')
        model.fit(X,y)
        print 'Intercepts:'
        print model.intercept_
        print 'Coefficients:'
        print model.coef_

```

Intercepts:

```
[-0.10153312 -0.56865552  0.67018864]
```

Coefficients:


```
[[ 2.80913084  1.10436996]
 [-1.8280927  -2.14320504]
 [-0.98103814  1.03883507]]
```

It is possible to get probabilities from the model using the method `predict_proba()`, but I have implemented the function `softmax_probs()` to evaluate probabilities with `model.intercept_` and `model.coef_` attributes.

```
In [5]: def softmax_probs(X, model):
        Xb = np.hstack((np.ones((X.shape[0],1)), X))
        thetas = np.hstack((model.intercept_.reshape((-1,1)), model.coef_))
        probs = np.exp(np.dot(Xb, np.transpose(thetas)))
        probs_sums = probs.sum(axis=1)
        probs = probs / probs_sums[:, np.newaxis]
        return probs

x1 = np.arange(-4.0, 3.0, 0.05)
x2 = np.arange(-3.0, 4.0, 0.05)
x1,x2 = np.meshgrid(x1, x2)
y_pred = softmax_probs(np.c_[x1.ravel(), x2.ravel()], model)
extent = -4.0, 3.0, -3.0, 4.0

fig = plt.figure(figsize=(10,6))
plt.imshow(y_pred[:,0].reshape(x1.shape), cmap=cm.Reds, alpha=.4,
           interpolation='bilinear',extent = extent, origin='lower')
plt.imshow(y_pred[:,1].reshape(x1.shape), cmap=cm.Blues, alpha=.4,
           interpolation='bilinear',extent = extent, origin='lower')
plt.imshow(y_pred[:,2].reshape(x1.shape), cmap=cm.Greens, alpha=.4,
           interpolation='bilinear',extent = extent, origin='lower')
plt.scatter(X[ix0,0],X[ix0, 1],marker='o',color='red')
plt.scatter(X[ix1,0],X[ix1, 1],marker='x',color='blue')
plt.scatter(X[ix2,0],X[ix2, 1],marker='s',color='green')
plt.show()
```




And another way to visualize it:

```
In [6]: fig = plt.figure(figsize=(16,8))
fig.add_subplot(131)
plt.imshow(y_pred[:,0].reshape(x1.shape), cmap=cm.Reds, alpha=.6,
           interpolation='bilinear', extent = extent, origin='lower')
plt.scatter(X[ix0,0],X[ix0, 1],marker='o',color='red')
plt.scatter(X[ix1,0],X[ix1, 1],marker='x',color='blue')
plt.scatter(X[ix2,0],X[ix2, 1],marker='s',color='green')


fig.add_subplot(132)
plt.imshow(y_pred[:,1].reshape(x1.shape), cmap=cm.Blues, alpha=.6,
           interpolation='bilinear', extent = extent, origin='lower')
plt.scatter(X[ix0,0],X[ix0, 1],marker='o',color='red')
plt.scatter(X[ix1,0],X[ix1, 1],marker='x',color='blue')
plt.scatter(X[ix2,0],X[ix2, 1],marker='s',color='green')

fig.add_subplot(133)
plt.imshow(y_pred[:,2].reshape(x1.shape), cmap=cm.Greens, alpha=.6,
           interpolation='bilinear', extent = extent, origin='lower')
plt.scatter(X[ix0,0],X[ix0, 1],marker='o',color='red')
plt.scatter(X[ix1,0],X[ix1, 1],marker='x',color='blue')
plt.scatter(X[ix2,0],X[ix2, 1],marker='s',color='green')

plt.show()
```



4 Generative learning algorithm



The logistic regression tries to find the straight line between two classes. This algorithm belongs to the class of **discriminative algorithms**. The idea of properties for discriminative algorithms is the following:

- learns $p(y \mid x)$
- or learns $h_{\theta}(x) \in \{0, 1\}$

Now we are going to talk about different approach. Assuming that we have two classes, we build two models: one - on positive examples only, the second - on negative examples only. When we get new sample we try to understand which model this sample match better and based on this matching we make a conclusion. This algorithm belongs to the class of **generative algorithms**. The idea of the generative algorithms is

- learn $p(x \mid y)$ and $p(y)$
- use Bayes rule to get $p(y = 1 \mid x) = \frac{p(x \mid y = 1)p(y)}{p(x)}$, where

$$p(x) = p(x \mid y = 0)p(y = 0) + p(x \mid y = 1)p(y = 1)$$